

Read Isolation, Consistency, and Recency

On this page

- [Isolation Guarantees](#)
- [Consistency Guarantees](#)
- [Recency](#)

Isolation Guarantees

Read Uncommitted

In MongoDB, clients can see the results of writes before the writes are [durable](#):

- Regardless of [write concern](#), other clients using ["local"](#) (i.e. the default) `readConcern` can see the result of a write operation before the write operation is acknowledged to the issuing client.
- Clients using ["local"](#) (i.e. the default) `readConcern` can read data which may be subsequently [rollback](#).

Read uncommitted is the default isolation level and applies to [mongod](#) standalone instances as well as to replica sets and sharded clusters.

Read Uncommitted And Single Document Atomicity

Write operations are atomic with respect to a single document; i.e. if a write is updating multiple fields in document, a reader will never see the document with only some of the fields updated.

With a standalone [mongod](#) instance, a set of read and write operations to a single document is serializable. With a replica set, a set of read and write operations to a single document is serializable *only* in the

absence of a rollback.

However, although the readers may not see a *partially* updated document, read uncommitted means the concurrent readers may still see the updated document before the changes are [durable](#).

Read Uncommitted And Multiple Document Write

When a single write operation modifies multiple documents, the modification of each document is atomic but the operation as a whole is not atomic and other operations may interleave. However, you can *isolate* single write operation that affects multiple documents using the [\\$isolated](#) operator.

Without isolating the multi-document write operations, MongoDB exhibits the following behavior:


1. Non-point-in-time read operations. Suppose a read operation begins at time t_1 and starts reading documents. A write operation then commits an update to one of the documents at some later time. The reader may see the updated version of the document, and therefore does not see a point-in-time snapshot of the data.
2. Non-serializable operations. Suppose a read operation reads a document d_1 at time t_1 and a write operation updates d_1 at some later time t_3 . This introduces a read-write dependency such that, if all operations were to be serialized, the read operation must precede the write operation. But also suppose that the write operation updates document d_2 at time t_2 and the read operation subsequently reads d_2 at some later time t_4 . This introduces a write-read dependency which would instead require the read operation to come *after* the write operation in a serializable schedule. There is a dependency cycle which makes serializability impossible.
3. Reads may miss matching documents that are updated during the course of the read operation.

Using the [\\$isolated](#) operator, a write operation that affects multiple documents can prevent other processes from interleaving once the write operation modifies the first document. This ensures that no reader sees the changes until the write operation completes or errors out.

[\\$isolated](#) does **not** work with [sharded clusters](#).

An isolated write operation does not provide “all-or-nothing” atomicity. That is, an error during the write operation does not roll back all its changes that preceded the error.

NOTE



`$isolated` operator causes write operations to acquire an exclusive lock on the collection, *even for document-level locking storage engines* such as WiredTiger. That is, `$isolated` operator will make WiredTiger single-threaded for the duration of the operation.

SEE ALSO

[Atomicity and Transactions](#)