

喧

原

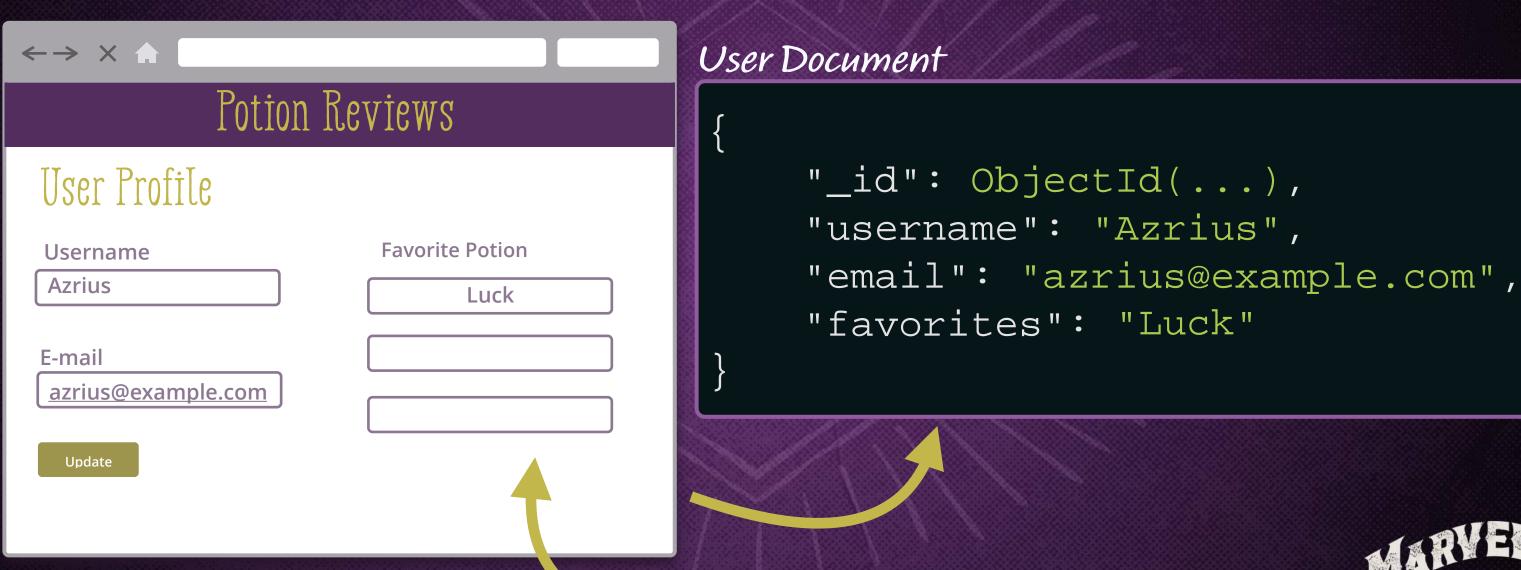
# Morphing Models

Level 4 - Section 1

**Data Modeling** 

#### Introducing User Profiles

We've added a user profile page to the site and want to allow users to enter up to 3 of their favorite potions.

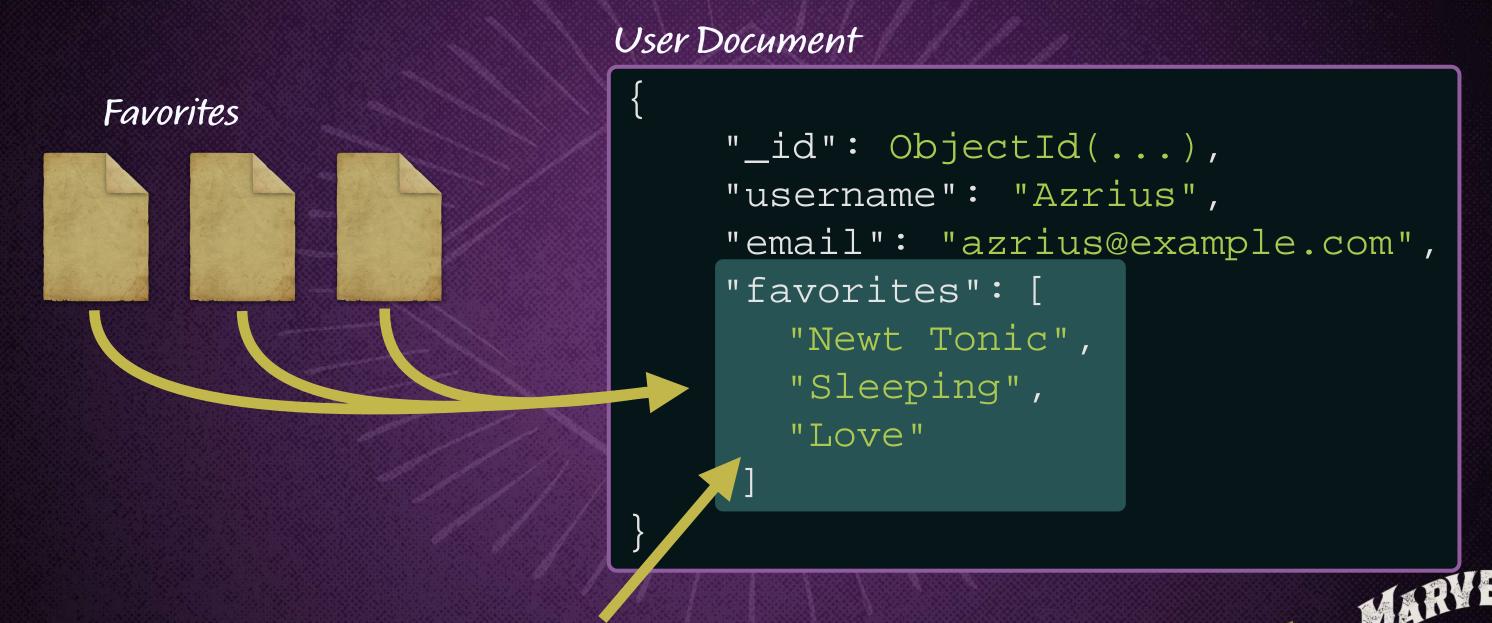


Need a way to store up to 3 potions



#### Storing Favorites Within a User

Users and their favorites are strongly related and will be used together often.



Use an array to hold multiple values

### **Adding Vendor Information**

We'd like to add more vendor information to our potions so users can be more informed about where they get their potions.

Potion Document

Vendor



Inserting the data as an embedded document

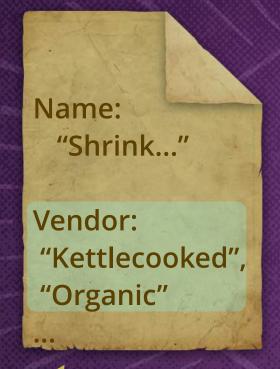
```
"_id": ObjectId(...),
"name": "Invisibility",
"vendor":
  "name": "Kettlecooked",
  "phone": 5555555555,
  "organic": true
```

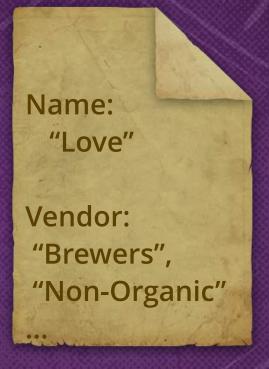


#### **Storing Vendor Information**

Each potion now contains its vendor information, which means we have vendor information repeated throughout our documents.











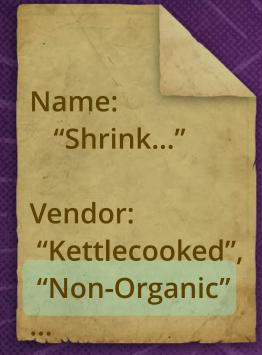




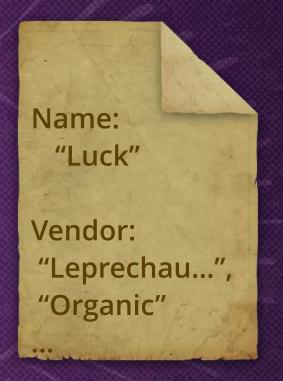
#### Dangers of Duplication

Duplicate data can be hard to keep consistent throughout the database.







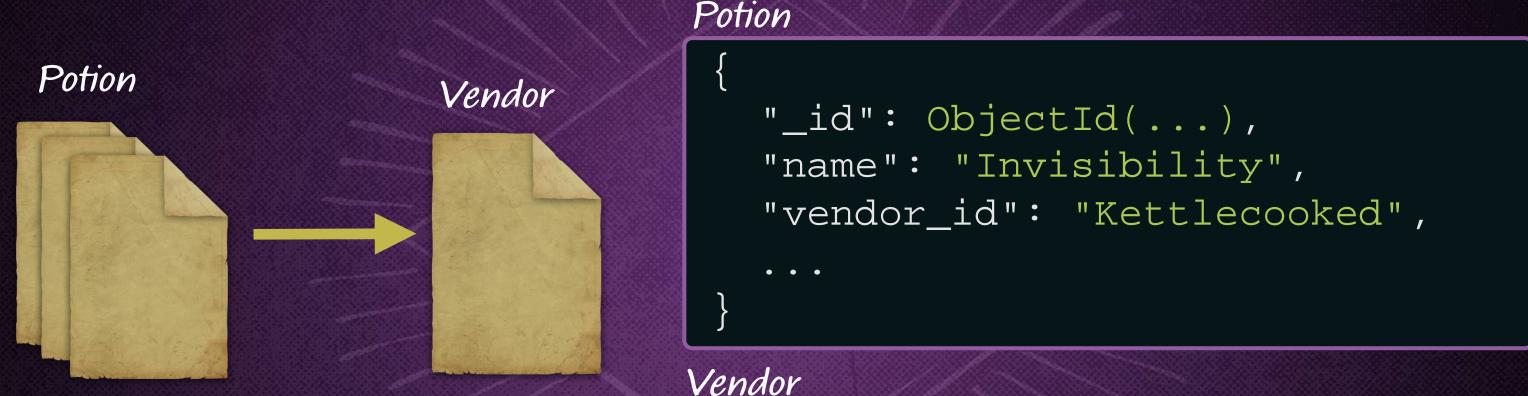


If 1 potion gets updated with new information and the rest don't, our data is no longer correct



#### Referencing Vendor Information

Instead of embedding the vendor information, we can create a vendors collection and reference the vendor document in each potion document.



Vendor names are unique and don't change

```
"_id": "Kettlecooked",
    "phone": 555555555,
    "organic": true
}
```

### **Inserting Referenced Documents**

```
> db.vendors.insert({
    "_id": "Kettlecooked",
    "phone": 555555555,
    "organic": true
})
```

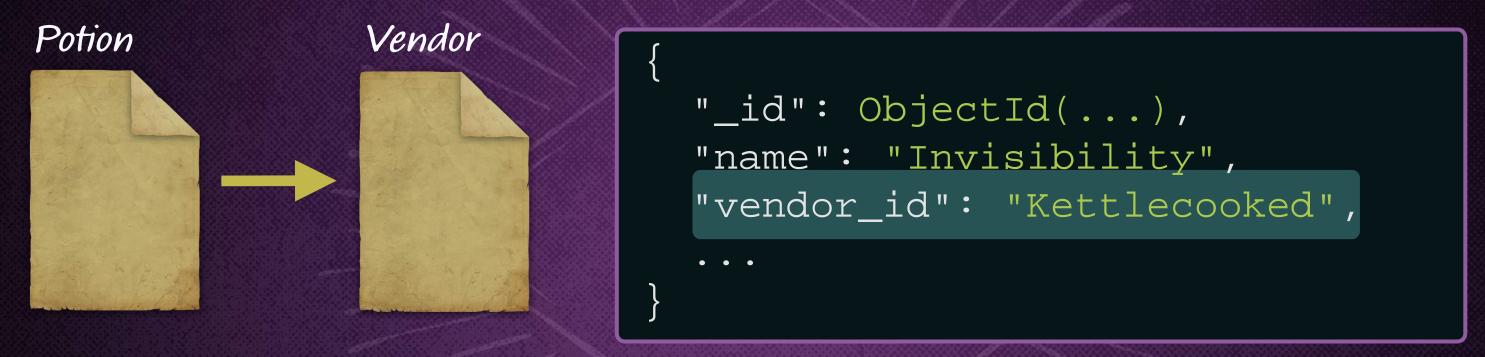
We can specify the unique \_id of document

```
> db.potions.insert({
    "name": "Invisibility",
    "vendor_id": "Kettlecooked"
    ...
})
```

Referenced document

#### Querying a Referenced Document

In order to pull a potion document and the vendor information, we must first query for the potion to get the vendor\_id and then query once more to get their vendor information.



First, query to retrieve potion information

```
db.potions.find({"name": "Invisibility"})

Second, query to retrieve vendor information

db.vendors.find({"_id": "Kettlecooked"})
```

We get the vendor \_id from the first query

#### Some Features of Embedded Documents

With a single query, we can grab a user's email and their addresses

SHELL



```
Data readily available
```

```
db.users.find({},{"email": true, "favorites": true})
```

```
{
    "_id": ObjectId(...),
    "email": "azrius@example.com"
    "favorites": [
        "Newt Tonic",
        "Sleeping",
        "Love"
    ]
}
```

#### **More Features of Embedded Documents**



#### Data readily available

db.users.find({},{"email": true, "favorites": true})

SHELI

```
Atomic write operations
```

SHELL

```
db.users.insert({
    "username": "Azrius",
    "email": "azrius@example.com",
    "favorites": ["Newt...", "Sleeping", "Love"]
})
```

Guarantee that the document write completely happens or doesn't at all

### Why Does Atomicity Matter?

If we update a user's email and add a favorite potion, but an error occurs in the favorites portion of the update, then none of the update will occur.

SHELL

SHELL

Updating email and adding a new favorite

db.users.update(...)

Document remains unchanged

Something bad happens

All fields updated!

db.users.update(...)

Successful WriteResult returned



### Referenced Documents Exist Independently

Since we've referenced vendor documents, they now exist entirely on their own outside of potion documents.





```
db.vendors.update({"_id": ObjectId(...)},{...})
```



Multi-document writes not supported

db.vendors.insert({...})

```
New Potion
SHELL

db.potions.insert({...})

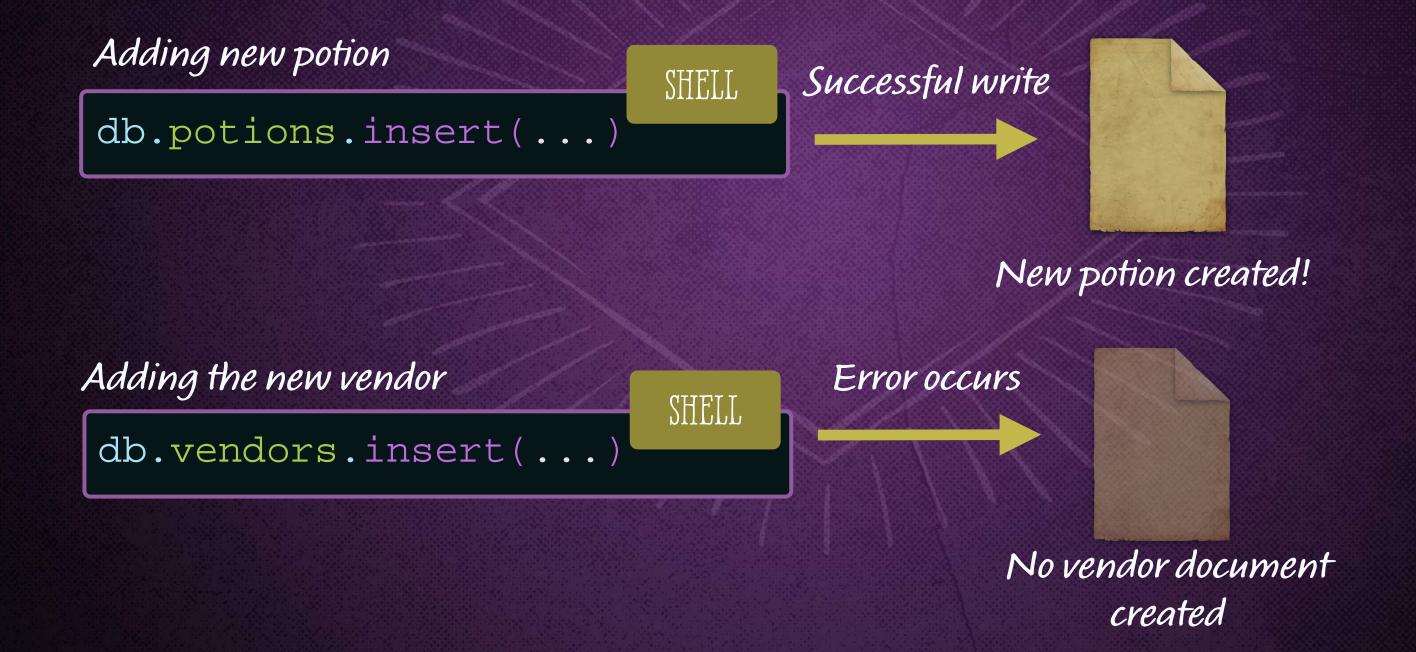
New Potion's Vendor
SHELL
```

No guarantee that both write operations will occur



#### Multi-document Write Operations

Currently, MongoDB doesn't support multi-document writes. We aren't able to guarantee write operations on an atomic level.



### Dangers of Not Having Transactions

MongoDB doesn't recognize document relationships, so we now have a potion with a nonexistent vendor.

Potion document



New potion created!



"vendor\_id": "Magical Inc.",

"\_id": ObjectId(...),

"name": "Time Travel",

Vendor never got created!

No vendor document created

The references still exist!

# Morphing Models

Level 4 – Section 2

**Data Modeling Decisions** 

### **Choosing Which Route to Take**

When deciding how to model a one-to-many relationship, we must carefully consider how our data will be used in order to decide between embedding and referencing.

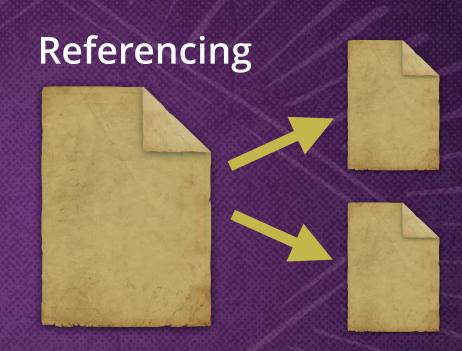




Single query

Documents accessed through parent

Atomic writes



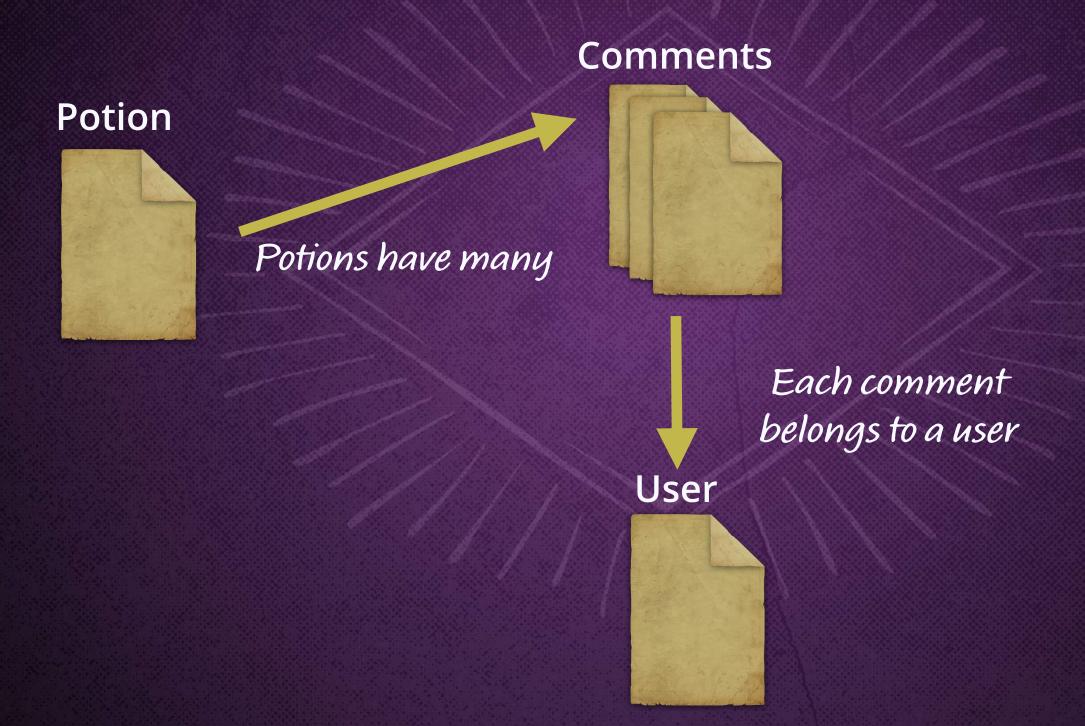
Requires 2 queries

Documents exist independently

Doesn't support multi-document writes

## Adding Comments to Potions

We'd like to allow users to comment on the potions and need to determine the best route.





#### First Question: How Will the Data Be Used?

Data that is frequently used together will benefit from being embedded while data that's rarely used can afford the cost of referencing.

How often is the data used together?	Always	Sometimes	Rarely
Embed			
Reference			

Embedding will work most of the time

Either option can work effectively

## Data Usage: Potion Comments



#### Potion Reviews



#### Invisibility

Score: 70

Taste: 4 Strength: 1

#### Comments

I don't agree. You are wrong.

-Sauron



X Referencing

We would be forced to perform multiple queries

Whenever we display potions, we'll always want to display comments

Username displayed for each comment



#### Second Question: What's the Size of the Data?

The size of the data in a relationship has a significant impact on data modeling. It's crucial to think ahead!

Expected Size	Less than 100	More than a few hundred	Thousands
Embed			
Reference			

Might start to see a decline in read performance when embedded

#### **Data Size: Potion Reviews**



#### Potion Reviews



#### Invisibility

Score: 70

Taste: 4 Strength: 1



I don't agree. You are wrong. -Sauron





When the data size is over 100, we can consider referencing

Most potions won't get more than 50 comments, and each comment has a single author



#### Third Question: Will the Data Change Often?

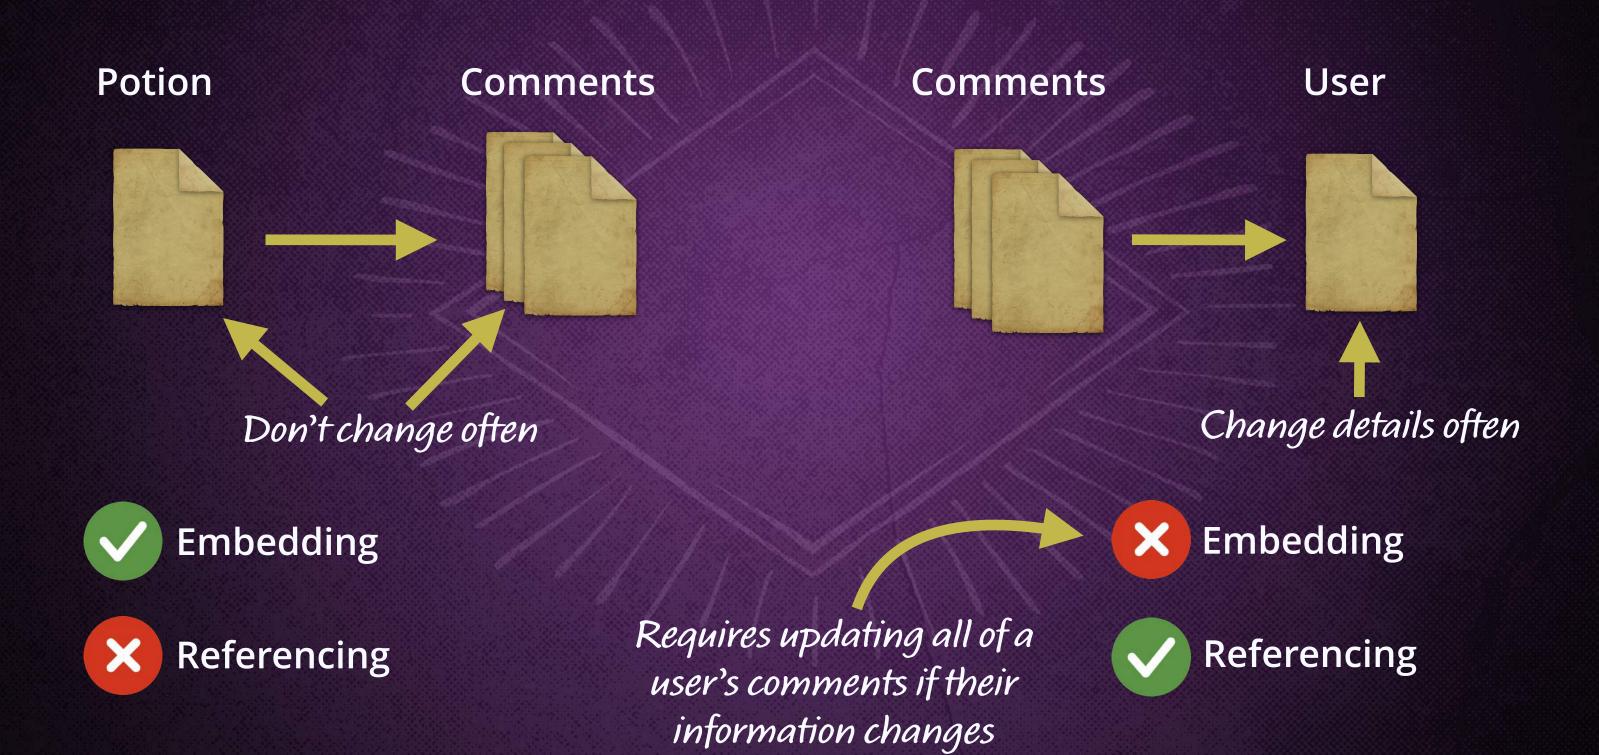
Sometimes embedding data can lead to data duplication. Depending on whether or not the duplicate data changes a lot can factor into our decision making.



Data duplication is okay if we don't expect change

Depends on whether or not we want the overhead of managing duplication

# Data Change: Potion Comments



#### **Embedding Comments in Potions**

We can confidently embed comments within potions. We know we'll have less than 100 comments per potion, they're used together often, and the data doesn't change often.

```
"name": "Invisibility",
"comments":[
  "title": "The best potion!",
  "body": "Lorem ipsum abra cadrbra"
```

Comments readily available

#### Referencing Users in Comments

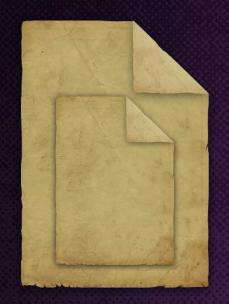
We only need the username from each user, and embedding user documents could lead to duplication issues. Let's reference a user by using his or her unique username.

```
"name": "Invisibility",
"comments":[
   "title": "The best potion!",
   "body": "Lorem ipsum abra cadrbra",
   "user_id": "Azrius"
```

Usernames can't be changed and are unique



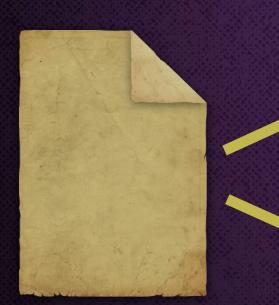
# Data Modeling Guidelines and Takeaways



Generally, embedding is the best starting point



If you find you need complex references, consider a relational database





Reference data when you need to access document independently



Consider referencing when you have large data sizes

